

Benjamin S. Manning

CONTACT	bmanning (at) mit.edu ◇ benjaminmanning.com	
EDUCATION	Massachusetts Institute of Technology Ph.D. Management Science, Information Technology Committee: <i>John Horton (Chair), Christopher Avery, Colin Camerer</i> S.M. Management Research	2022 - Present 2024
	Harvard University M.P.P. Economic & Social Policy	2021
	Washington University in St. Louis B.A. Mathematics (College Honors in Arts & Sciences)	2017
GRANTS (\$488,325)	Google Cloud Research Credits Grant (\$5,000; w/ Alex Moehring) Schmidt Sciences AI at Work Award (\$193,325) Emergent Ventures Fast Grant (\$50,000) MIT GenAI Impact Consortium & OpenAI (\$180,000; w/ John J. Horton) Schmidt Sciences Grant Proposal Award (\$10,000) Emergent Ventures Fast Grant (\$50,000)	2026 2025 2025 2025 2025 2024
AWARDS & SCHOLARSHIPS	EC '26 Exemplary Paper Award WISE Best Paper Award Harvard Distinction in Student Teaching Award Becker Friedman Institute Price Theory Summer Camp Kosciuszko Foundation Tuition Scholarship (\$3,500) Kosciuszko Foundation Tuition Scholarship (\$3,000)	2026 2025 2025 2025 2022 2019
JOB MARKET PAPER	General Social Agents (Benjamin S. Manning ¹ & John J. Horton) Revise & Resubmit at <i>Econometrica</i> • WISE 2025 Best Paper Award • Extended abstract at <i>ACM Conference on Economics & Computation (EC '26)</i> • Media: Marginal Revolution	
	Abstract: Useful social science theories predict behavior across settings. However, applying a theory to make predictions in new settings is challenging: rarely can it be done without ad hoc modifications to account for setting-specific factors. We argue that AI agents put in simulations of those novel settings offer an alternative for applying theory, requiring minimal or no modifications. We present an approach for building such “general” agents that use theory-grounded natural language instructions, existing empirical data, and knowledge acquired by the underlying AI during training. To demonstrate the approach in settings where no data from that data-generating process exists—as is often the case in applied prediction problems—we design a heterogeneous population of 883,200 novel games. AI agents are constructed using human data from a small set of conceptually related but structurally distinct “seed” games. In preregistered experiments, on average, agents predict initial human play in a random sample of 1,500 games from the population better than (i) a cognitive hierarchy model, (ii) game-theoretic equilibria, and (iii) out-of-the-box agents. For a small set of separate novel games, these simulations predict responses from a new sample of human subjects <i>better</i> even than the most plausibly relevant published human data.	
WORKING PAPERS	Automated Social Science: Language Models as Scientist and Subjects (Benjamin S. Manning ¹ , Kehang Zhu ¹ , & John J. Horton) Reject & Resubmit at <i>The Quarterly Journal of Economics</i> • EC '26 Exemplary Paper Award (one of six)	

¹Indicates first author or co-first author

- Extended abstract at *ACM Conference on Economics & Computation (EC '26)*
- Media: [Marginal Revolution](#)

Abstract: We present an approach for automatically generating and testing, *in silico*, social scientific hypotheses. This automation is made possible by recent advances in large language models (LLM), but the key feature of the approach is the use of structural causal models. Structural causal models provide a language to state hypotheses, a blueprint for constructing LLM-based agents, an experimental design, and a plan for data analysis. The fitted structural causal model becomes an object available for prediction or the planning of follow-on experiments. We demonstrate the approach with several scenarios: a negotiation, a bail hearing, a job interview, and an auction. In each case, causal relationships are both proposed and tested by the system, finding evidence for some and not others. We provide evidence that the insights from these simulations of social interactions are not available to the LLM purely through direct elicitation. When given its proposed structural causal model for each scenario, the LLM is good at predicting the signs of estimated effects, but it cannot reliably predict the magnitudes of those estimates. In the auction experiment, the *in silico* simulation results closely match the predictions of auction theory, but elicited predictions of the clearing prices from the LLM are inaccurate. However, the LLM’s predictions are dramatically improved if the model can condition on the fitted structural causal model. In short, the LLM knows more than it can (immediately) tell.

Modeling Human Behavior with Type Vectors Using AI

(Matthew O. Jackson¹, Benjamin S. Manning¹, Yutong Xie¹, Walter Yuan, & Qiaozhu Mei)

Under Review

Abstract: We introduce a general, easy-to-implement AI-based method for modeling and analyzing the structure and complexity of human behavior. We assign a large language model a “type vector” and then prompt it to choose actions across settings in which we observe human choices. For instance, the type vector (2, 4) becomes “You are a player characterized by the following profile: Altruism: 2 out of 5, Risk Aversion: 4 out of 5,” after which it is prompted to make choices. We vary the dimensions (e.g., Altruism, Fairness, Trust, ...) and values (e.g., 1–5) to minimize distance to human choices. Applying the method to 119,147 decisions made by 78,657 subjects from more than 35 countries across 10 classic economic game roles, we find that human behavior can be closely matched using three dimensions: Risk Aversion, Strategic Sophistication, and Trust. Moreover, the types needed to fit individuals across games cluster into fewer than a dozen groups, and can predict behavior in held-out games with different rules and available actions. The results suggest that behavior across diverse settings can be approximated by a low-dimensional, portable representation, supporting the possibility of general yet parsimonious theories across the behavioral sciences. More broadly, this new modeling method can provide insights into the structure of many human behaviors.

Self-Supervised Preference Learning for Multimodal Foundation Models

(Akshata Tiwari¹, Jillian Ross, Benjamin S. Manning, & Andrew Lo)

Under Review

Preliminary Abstract (subject to change): Preference optimization for multimodal foundation models typically draws its signal from external judgment: human annotations, model-based scoring, or task-specific supervision. We propose a self-supervised alternative based on a simple principle: if two images depict similar underlying data, their descriptions should be similar. We instantiate this idea for time-series data rendered as charts—a setting that is simple, broadly important, and well-suited to clean downstream tasks. Two charts are neighbors if they trace similar patterns in the underlying data; we compute these neighborhoods directly from the data, generate candidate descriptions of each image and its neighbors using the model itself, and rank candidates by their agreement with the neighborhood to form preference pairs. We apply this method to direct preference optimization on two open-source foundation models, Qwen2-VL-7B and Gemma-3-4B, across three domains: financial price paths, electrocardiograms, and seismic traces. The learned signal transfers zero-shot to held-out tasks. Relative to the supervised fine-tuning baseline, the detection of large price-moves increases by 31%, crash-risk prediction by 14%, electrocardiogram abnormality detection by 10%, and seismic event detection by 3.3%. Probes show the gains reflect stronger alignment between visual inputs and textual outputs, and ablations confirm that neither generic preference optimization nor learned-feature neighborhoods produce the same gains.

SELECTED
WORK IN
PROGRESS

When Stated, Revealed, and Normative Preferences Diverge: Experimental Evidence from Personalized News Feeds
(Benjamin S. Manning¹ & Alex Moehring¹)

Taste-as-Program: Scalable Preference Elicitation with AI

(Gili Rusak¹, Benjamin S. Manning¹, & John J. Horton)

Mass In Silico Replication of Social Science Experiments

(Jacob Snyder¹, Benjamin S. Manning¹, John J. Horton¹, & Abhishek Nagara¹)

What Causes Price Bubbles? Theory and Evidence from Simulations and the Lab

(Cady Ngo¹, Benjamin S. Manning¹, Christopher Avery¹, & Colin Camerer¹)

Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus?

(John J. Horton¹, Apostolos Filippas¹, Benjamin S. Manning¹)

The Review of Economics and Statistics, Conditionally Accepted

- Extended abstract at *ACM Conference on Economics & Computation (EC '24)*

Abstract: We argue that newly-developed large language models (LLMs), because of how they are trained and designed, are implicit computational models of humans—a *Homo silicus*. LLMs can be used like economists use *Homo economicus*: they can be given endowments, information, preferences, and so on, and then their behavior can be explored in scenarios via simulation. Experiments using this approach, derived from Charness and Rabin (2002), Kahneman et al. (1986), Samuelson and Zeckhauser (1988), Oprea (2024b), and Horton (2025), show qualitatively similar results to the original, and when they differ, it is often generative for future research. We discuss potential applications, conceptual issues, and why this approach can inform the study of humans.

Prompt Adaptation as a Dynamic Complement in Generative AI Systems

(Eaman Jahani¹, Benjamin S. Manning, Joe Zhang, Hong-Yi TuYe, Mohammed Alsobay, Christos Nicolaides, Siddharth Suri, & David Holtz)

Information Systems Research, 2026

- Media: [Marginal Revolution](#), [MIT Ideas Made to Matter](#), [Columbia Business School Insights](#)

Abstract: As generative AI systems rapidly improve, a key question emerges: how do users adapt to these changes, and when does such adaptation matter for realizing performance gains? Drawing on theories of dynamic capabilities and IT complements, we study *prompt adaptation*—how users adjust their inputs in response to evolving model behavior—using a common experimental design applied to two preregistered tasks with 3,750 total participants who submitted nearly 37,000 prompts. We show that the importance of prompt adaptation depends critically on task structure. In a task with fixed evaluation criteria and an unambiguous goal, user prompt adaptation accounts for roughly half of the performance gains from a model upgrade. In contrast, in an open-ended creative task where the space of acceptable outputs is effectively unbounded and quality is subjective, performance improvements are driven primarily by model capability; prompt adaptation plays a limited role. We further show that automated prompt rewriting cannot generally substitute for human adaptation: when aligned with task objectives, it can modestly improve performance, but when misaligned, it can actively undermine the gains from model improvements. Together, these findings position prompt adaptation as a dynamic complement whose importance depends on task structure and system design, and suggest that without it, a substantial share of the economic value created by advances in generative models may go unrealized.

National Megastudy Shows that Email Nudges to Elementary School Teachers Boost Student Math Achievement, Particularly When Personalized

(Angela Duckworth¹, Katherine Milkman, ... Benjamin S. Manning, ..., & 26 others)

Proceedings of the National Academy of Sciences, 2025

Abstract: In response to the alarming recent decline in US math achievement, we conducted a national megastudy in which 140,461 elementary school teachers who collectively taught 2,992,027 students were randomly assigned to receive a variety of behaviorally informed email nudges aimed at improving students' progress in math. Specifically, we partnered with the nonprofit educational platform Zearn Math to compare the impact of 15 different interventions with a reminder-only megastudy control condition. All 16 conditions entailed weekly emails delivered to teachers over 4-wk in the fall of 2021. The best-performing intervention, which encouraged teachers to log into Zearn Math for an updated report on how their students were doing that week, produced a 5.06% increase in students' math progress (3.30% after accounting for the winner's curse). In exploratory analyses, teachers who received any behaviorally informed email nudge (vs. a reminder-only megastudy control) saw their students' math progress boosted by an average of 1.89% during the 4-wk intervention period; emails referencing personalized data (i.e., classroom-specific statistics) outperformed emails that did not by 2.26%. While small in size, these intervention effects were consistent across school socioeconomic status and school type (public, private, etc.) and, further, persisted in the 8-wk

post-intervention period. Collectively, these findings underscore both how difficult it is to change behavior and the need for large-scale, rigorous, empirical research of the sort undertaken in this megastudy.

Effect Size Magnification: No Variable is as Important as the One You're Thinking About—While You're Thinking About It

(Linnea Gandhi¹, **Benjamin S. Manning**, & Angela Duckworth)

Current Directions in Psychological Science, 2024

- Top 10 most downloaded Association for Psychological Science (APS) journal articles in 2024
- Media: **Under the Cortex**

Abstract: The goal of psychological science is to discover truths about human nature, and the typical form of empirical insights is a simple statement of the form x relates to y . We suggest that such “one-liners” imply much larger x - y relationships than those we typically study. Given the multitude of factors that compete and interact to influence any human outcome, small effect sizes should not surprise us. And yet they do—as evidenced by the persistent and systematic underpowering of research studies in psychological science. We suggest an explanation. Effect size magnification is the tendency to exaggerate the importance of the variable under investigation because of the momentary neglect of others. Although problematic, this attentional focus serves a purpose akin to that of the eye's fovea. We see a particular x - y relationship with greater acuity when it is the center of our attention. Debiasing remedies are not straightforward, but we recommend (a) recalibrating expectations about the effect sizes we study, (b) proactively exploring moderators and boundary conditions, and (c) periodically toggling our focus from the x variable we happen to study to the non- x variables we do not.

REFEREED CONFERENCE PROCEEDINGS

Strategic Tradeoffs Between Humans and AI in Multi-Agent Bargaining

(Crystal Qian¹, Kehang Zhu¹, John J. Horton, **Benjamin S. Manning**, Vivian Tsai, James Wexler, & Nithum Thain)

ACM Conference on Intelligent User Interfaces (IUI), 2026

Abstract: Markets increasingly accommodate large language models (LLMs) as autonomous decision-making agents. As this transition occurs, it becomes critical to evaluate how these agents behave relative to their human and task-specific statistical predecessors. In this work, we present results from an empirical study comparing humans ($N=216$), multiple frontier LLMs, and customized Bayesian agents in dynamic multi-player bargaining games under identical conditions. Bayesian agents extract the highest surplus with aggressive trade proposals that are frequently rejected. Humans and LLMs achieve comparable aggregate surplus within their groups, but exhibit different trading strategies. LLMs favor conservative, concessionary proposals that are usually accepted by other LLMs, while humans propose trades that are consistent with fairness norms but are more likely to be rejected. These findings highlight that performance parity—a common benchmark in agent evaluation—can mask substantive procedural differences in how LLMs behave in complex multi-agent interactions.

BOOK CHAPTERS

The Coasean Singularity? Demand, Supply, and Market Design with AI Agents

(Peyman Shahidi¹, Gili Rusak¹, **Benjamin S. Manning**¹, Andrey Fradkin¹, & John J. Horton¹)

Chapter 6 in *The Economics of Transformative AI*, University of Chicago Press, 2026

- Media: **Marginal Revolution**, MIT Initiative on the Digital Economy

Abstract: AI agents—autonomous systems that perceive, reason, and act on behalf of human principals—are poised to transform digital markets by dramatically reducing transaction costs. This chapter evaluates the economic implications of this transition, adopting a consumer-oriented view of agents as market participants that can search, negotiate, and transact directly. From the demand side, agent adoption reflects derived demand: users trade off decision quality against effort reduction, with outcomes mediated by agent capability and task context. On the supply side, firms will design, integrate, and monetize agents, with outcomes hinging on whether agents operate within or across platforms. At the market level, agents create efficiency gains from lower search, communication, and contracting costs, but also introduce frictions such as congestion and price obfuscation. By lowering the costs of preference elicitation, contract enforcement, and identity verification, agents expand the feasible set of market designs but also raise novel regulatory challenges. While the net welfare effects remain an empirical question, the rapid onset of AI-mediated transactions presents a unique opportunity for economic research to inform real-world policy and market design.

TEACHING

BUSN 38802: Managerial Decision Making; **TA for Richard Thaler**, University of

Chicago—Booth London Campus, Fall 2024 (Rating: 4.7/5.0)

API-101Z: Markets and Failures; **TA for Christopher Avery**—Harvard University, Fall 2024 (Rating: 4.6/5.0)

API-101Z: Markets and Failures; **TA for Christopher Avery**—Harvard University, Fall 2023 (Rating: 4.3/5.0)

GUEST
LECTURES

API-101Z: Markets and Failures; **Christopher Avery**—Harvard University, Fall 2025

QSTBA 830: Business Experimentation and Causal Methods; **Andrey Fradkin**—Boston University, Spring 2025

AFFILIATIONS

MIT Initiative on the Digital Economy, Heterodox Academy

REFEREEING

Economics: *Quarterly Journal of Economics*, *Experimental Economics*, *Southern Economic Journal*

General Interest: *Proceedings of the National Academy of Sciences*, *Proceedings of the National Academy of Sciences Nexus*

Management: *Management Science*, *Manufacturing and Service Operations Management (MSOM)*, *International Conference on Information Systems (ICIS)*

Computer Science: *Conference on Neural Information Processing Systems (NeurIPS)*, *ACM Conference on Collective Intelligence*

Sociology: *Sociological Science*

Program Committees: *ACM Conference on Economics and Computation (2026)*, *The Business Implications of Generative AI (BIG.AI) (2026)*, *Workshop on Agent Behavior at the Conference on Language Modeling (COLM) (2026)*, *Conference on Digital Experimentation (CODE) (2023)*

Grants: *J-PAL Project AI Evidence*

INVITED
SEMINARS

2026 / 2027 (Includes scheduled talks): University of North Dakota (Real Estate, Economics & Finance); LSE (Centre for Philosophy of Natural and Social Science); University of Connecticut (Operations & Information Management); Chapman University (Business and Economics)

2025: Dartmouth College (Economics); Instacart Economics Team

2024: Instacart Economics Team

CONFERENCE
PRESENTATIONS

2026 / 2027 (Includes scheduled talks): Zurich Workshop on Economics and Psychology; Informs Annual Meeting; Conference on Information Systems and Technology (CIST); AI in Social Science Conference (UChicago); AI and Economics Summer Conference (UChicago); NBER AI and Economic Measurement; Columbia Management, Analytics, and Data Conference; Wharton AI and the Future of Work Conference; Digital Economy Workshop (Athens); AEA Annual Meeting.

2025: Workshop on Information Systems and Economics (WISE); North American Meeting of the Economic Science Association (ESA); KOF ETH Zurich and IZA Workshop: Matching Workers and Jobs Online; Machine Learning in Economics Summer Conference (UChicago); Academy of Management (AOM) Annual Meeting; Wharton AI and the Future of Work Conference; NBER/CEME Decentralization and AI; Khipu Latin American Meeting in AI; AEA Annual Meeting.

2024: National Association for Business Economics (NABE); Econometric Society Conference on Economics and AI and ML; ZEW Conference on the Economics of ICT; International Conference of the French Association of Experimental Economics; Statistical Conference in E-Commerce Research (SCECR); Wharton AI and the Future of Work Conference; Microsoft Research AI, Cognition, and the Economy Workshop; Measuring Development: AI, the Next Generation at the World Bank.

2023: MIT Computer Science and AI Lab (CSAIL) (FutureTech)

EMPLOYMENT	University of Pennsylvania	2021 - 2022
	Research Assistant for Angela Duckworth	
	University of Chicago	2020 - 2022
	Research Assistant for Jon Rogowski	
	Dartmouth College	2018 - 2019
	Research Assistant for Sydney Finkelstein	
	Phetpittayakom School	2017 - 2018
	High School Math Teacher in Thailand	
	Optiver US LLC	2016
	Execution Trading Intern	